Mainstream: Adaptive compute sharing for video analysis

Angela Jiang, Daniel Wong, Christopher Canel, Ishan Misra, Michael Kaminsky, Michael Kozuch, Padmanabhan Pillai, David Andersen, Greg Ganger

Overview

Goal:

- Efficiently run concurrent streaming video analysis apps

Problem:

- Most video analysis apps perform DNN inference
  - Running several full DNNs becomes very slow

Mainstream:

- Identifies and shares redundant DNN computation
  - By exploiting nature of fine-tuned DNNs
- Decides at runtime how much to share
  - Balances specialization vs. sharing trade-off
- Optimizes when hardware and set of apps is known

Transfer Learning

- When training task B, use DNN pre-trained for task A
  - Common practice for training networks

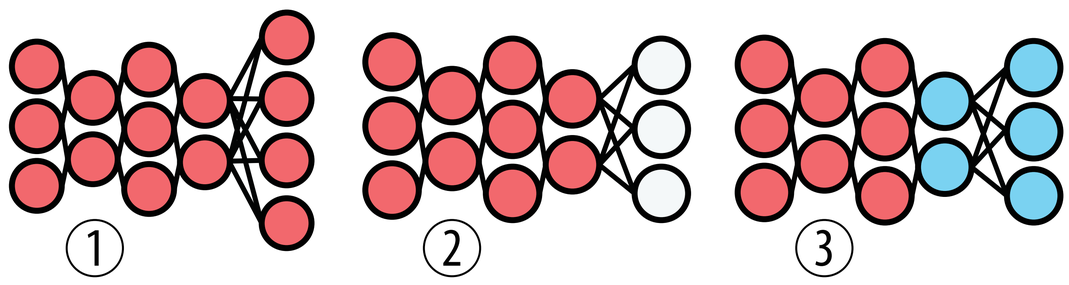

1. Network is trained from scratch for task A (e.g., ImageNet)
2. Replace A-specific final layer with B-specific final layer
3. Fine-tune part of network for task B, other layers held frozen

Sharing Computation

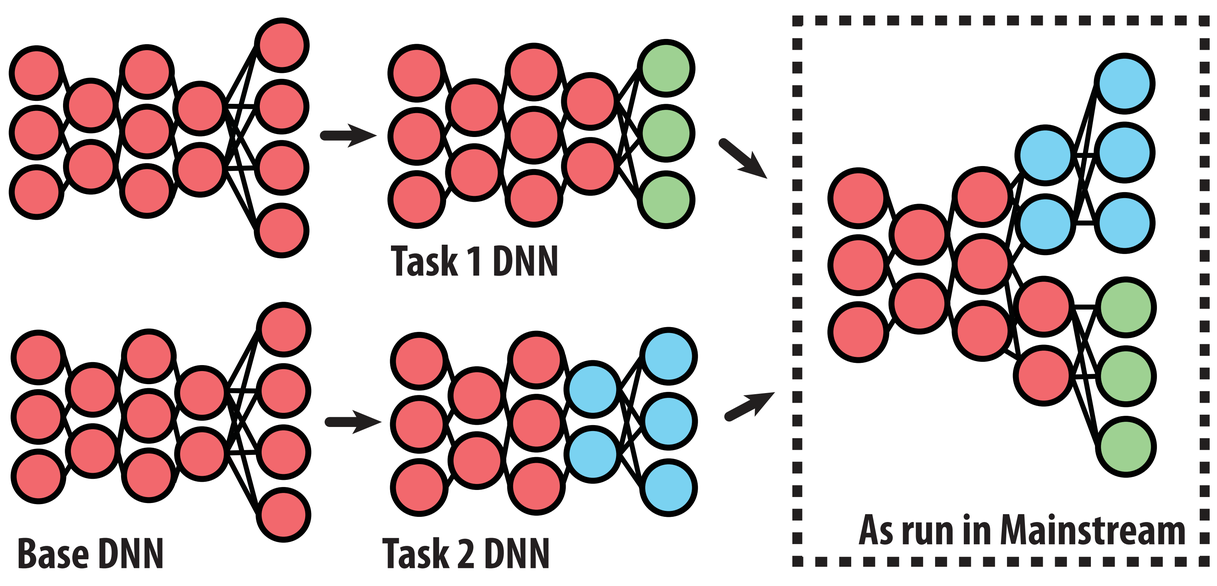

Mainstream Architecture

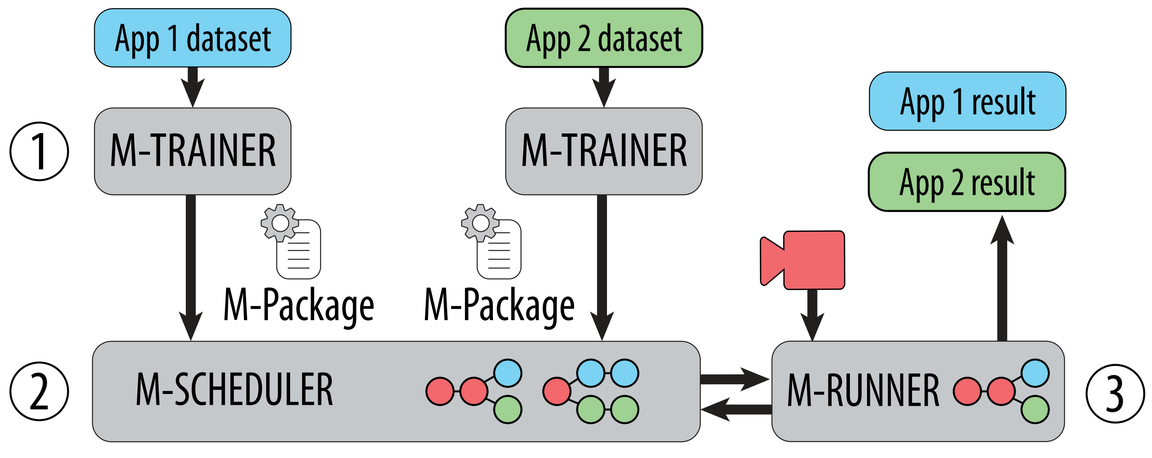

1. **M-trainer** trains DNNs with varying % of network held frozen
2. **M-Scheduler** determines amount of DNN to share for each app
3. **M-Runner** processes video stream using deployed DNNs

Specialization vs. Sharing Trade-off

Experimental setup: Train image classifiers to recognize flowers. Run simultaneous classification pipelines on an Intel NUC.

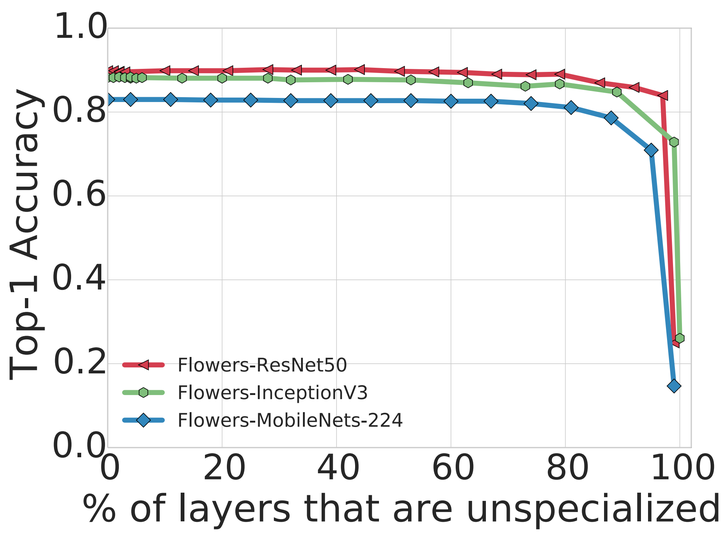

Less specialization → Lower per-frame acc.

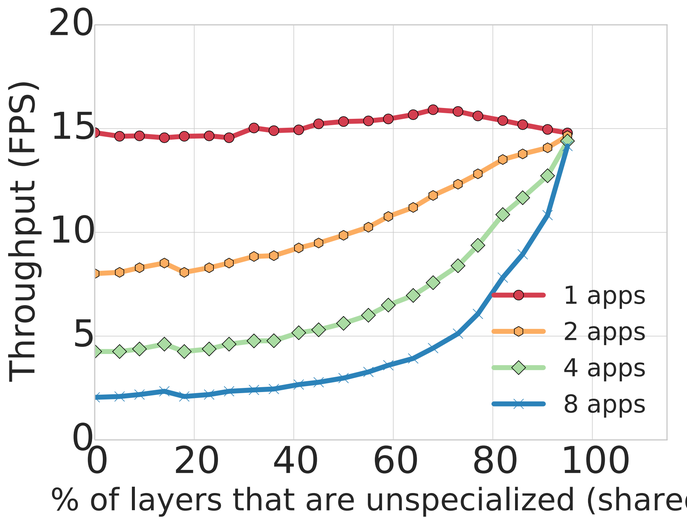

Less specialization → Higher throughput

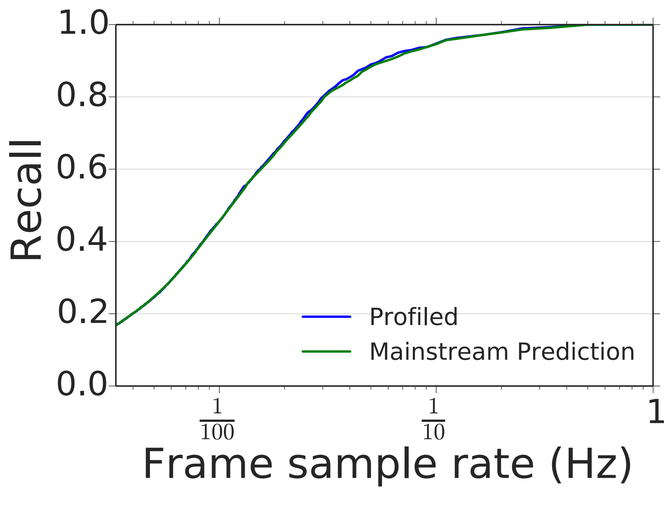

Use analytical model to find navigate trade-offs

Application Performance

Recall: % of events detected; **Precision:** % of detected events that are correct; **F1 score:** Harmonic mean of precision and recall

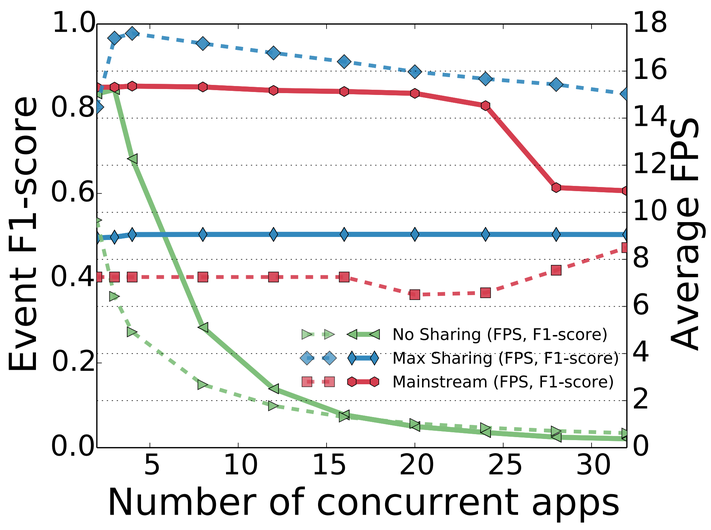

- "No Sharing" deploys full DNN for each app
- "Max Sharing" shares all but final layer
- Mainstream achieves up to 28X higher F1

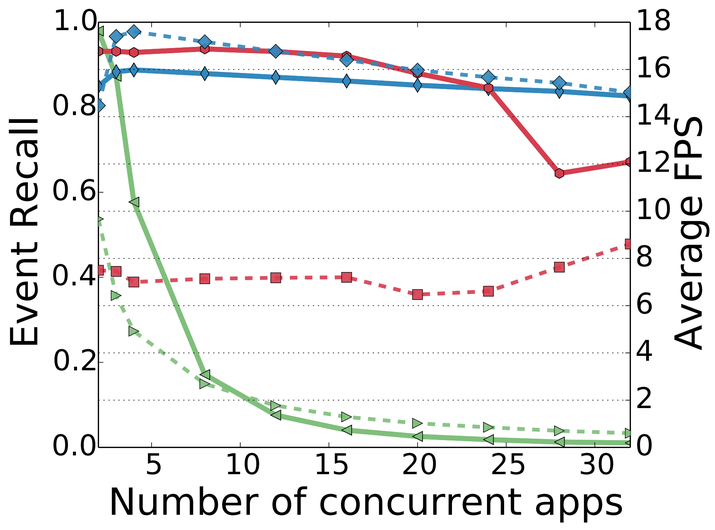

- "No Sharing" (NS) has low FPS, high acc
- NS misses events, incurring low recall
- Mainstream balances FPS and acc

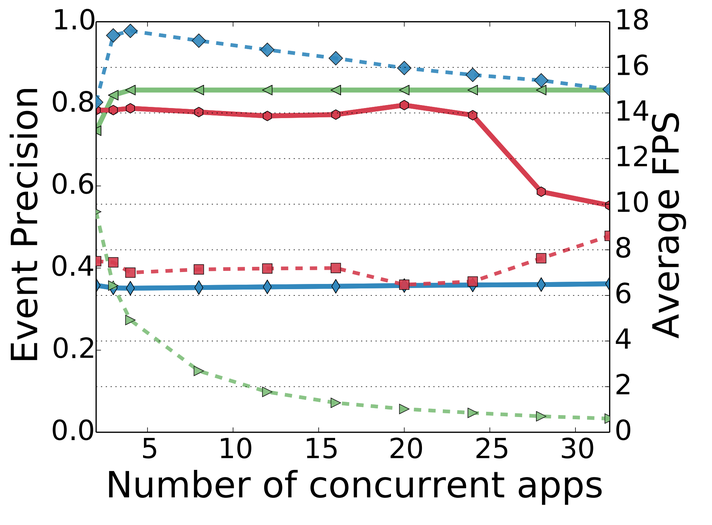

- "Max Sharing" (Max) has high FPS, low acc
- Max's false positives cause low precision
- Mainstream balances precision and recall